

## Authentication

<b>Service Account tokens</b>	<b>3</b>
<b>Client certs</b>	<b>3</b>
<b>OpenID connect tokens</b>	<b>3</b>
<b>Bootstrap Tokens - Beta</b>	<b>3</b>
<b>Static token file</b>	<b>3</b>
<b>Static password file</b>	<b>4</b>
<b>Security Context</b>	<b>4</b>
runAsUser=405	5
runAsGroup	5
runAsNonRoot=true	5
privileged=true	5
Linux Capabilities	7
ReadOnlyRootFilesystem	7
SELinux	8
<b>Secure Images</b>	<b>8</b>
Access private repositories	9
<b>Authorization</b>	<b>9</b>
Modes	9
Node - A special-purpose authorization mode that grants permissions to kubelets based on the pods they are scheduled to run	9
RBAC -	9
Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within an enterprise	9
Webhook -	9
ABAC - Attribute-based access control (ABAC) defines an access control paradigm whereby access rights are granted to users through the use of policies which combine attributes together. The policies can use any type of attributes (user attributes, resource attributes, object, environment attributes, etc).	10
<b>Network Policies</b>	<b>10</b>
<b>Creating TLS certificates</b>	<b>10</b>
<b>Admission Controller</b>	<b>10</b>
AlwaysPullImages	10

You can enable multiple authentication methods at once. You should usually use at least two methods:

- **service account tokens** for service accounts - Enabled by default
- at least one other method for **user authentication**.
  - Client certs
  - OpenID connect tokens
  - Bootstrap Tokens - Beta
  - Static token file
  - Static password file

users/groups are not represented as objects in kubernetes

### Serviceaccount

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

<https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/>

- represented as object
- Used by pods/applications to access API servers
- Have associated **secrets and generated automatically and mounted automatically to POD**

**To disable → automountServiceAccountToken: false**

**In service account yaml or POD yaml..**

- **The service account has to exist at the time the pod is created, or it will be rejected.**
- **You cannot update the service account of an already created pod.**
- **Any tokens for non-existent service accounts will be cleaned up by the token controller.**
- Secrets contains
  - Jwt TOKEN -- **Created and signed by using the service account key pair**
    - **service-account-key.pem**
    - **Service-account.pem**
    - Controller manager generated tokens..
  - Service account cert (service-account.pem)
    - Required for validating the JWT token...

### A Service account admission controller

1. - adding default service account to every pod if not present
2. - mounting secret as volume

**Token controller -**

1. creates new token/secrets for new service account..Also takes care of deletion..

#### **Service account controller-**

1. ensures every new namespace have a default service account

## Service Account tokens

- Enabled by default
- Service account needs to pass signed tokens to authenticate
- Request authenticated using service-account keys
- While starting API server below param needs to be passed :  
**--service-account-key-file** A file containing a PEM encoded key for signing bearer tokens. **If unspecified, the API server's TLS private key will be used.**

**Service accounts authenticate with the username `system:serviceaccount:(NAMESPACE):(SERVICEACCOUNT)`, and are assigned to the groups `system:serviceaccounts` and `system:serviceaccounts:(NAMESPACE)`.**

## Client certs

- enabled by passing the **--client-ca-file=SOMEFILE** option to API server
- SOMEfile - CA certs that will validate the client certs. There can be one or more CA certs file specified when starting API server.
- Hence client need to specify which authority granted the cert...
- Client cert -- > CN field == UserName  
O field == Group name, user can be part of more than one group..  
`openssl req -new -key jbeda.pem -out jbeda-csr.pem -subj "/CN=jbeda/O=app1/O=app2"`

## OpenID connect tokens

## Bootstrap Tokens - Beta

## Static token file

- Enabled by passing **--token-auth-file=SOMEFILE** option to API server.
- SOMEFILE - static CSV file containing below contents..
  - **token,user,uid,"group1,group2,group3"**
- Token file edit - needs API server restart
- Tokens do not expire, last indefinitely..
- Add the token in HTTP header

◆ Authorization: Bearer 31ada4fd-adec-460c-809a-9e56ceb75269

## Static password file

- Enabled by passing **--basic-auth-file=SOMEFILE** option to API server.
  - SOMEFILE - static CSV file containing below contents..
    - **token,user,uid,"group1,group2,group3"**
  - password edit - needs API server restart
  - password do not expire, last indefinitely..
  - Add the token in HTTP header
- ◆ Authorization: Basic **BASE64ENCODED(USER:PASSWORD)**

## Security Context

1. A security context defines privilege and access control settings for a Pod or Container
2. Security settings that you specify for a Container apply only to the individual Container, and they override settings made at the Pod level when there is overlap.

apiVersion: v1

kind: Pod

metadata:

name: security-context-demo

spec:

securityContext:

runAsUser: 1000

runAsGroup: 3000

fsGroup: 2000

volumes:

- name: sec-ctx-vol

emptyDir: {}

containers:

- name: sec-ctx-demo

image: busybox

command: [ "sh", "-c", "sleep 1h" ]

volumeMounts:

- name: sec-ctx-vol

mountPath: /data/demo

securityContext:

allowPrivilegeEscalation: **false**

## runAsUser=405

- The UID to run the entrypoint of the container process. Defaults to user specified in image metadata if unspecified. May also be set in PodSecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.

## runAsGroup

- The GID to run the entrypoint of the container process. Uses runtime default if unset. May also be set in PodSecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.
- The runAsGroup field specifies the primary group ID of 3000 for all processes within any containers of the Pod. If this field is omitted, the primary group ID of the containers will be root(0).

## runAsNonRoot=true

- Indicates that the container must run as a non-root user.
- If true, the Kubelet will validate the image at runtime to ensure that it does not run as UID 0 (root) and fail to start the container if it does. If unset or false, no such validation will be performed. May also be set in PodSecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.
- Error:
  - **Error: container has runAsNonRoot and image will run as root**

## privileged=true

- Run container in privileged mode. Processes in privileged containers are essentially equivalent to root on the host. Defaults to false.
- Pods default ---- > has root privileges only on the container process

```
imsrv01@imsrv01-Lenovo-ideapad-700-15ISK:~/CKA/cka/security$ kubectl exec -it pod-with-defaults -- ls /dev
```

core	null	shm	termination-log
fd	ptmx	stderr	tty
full	pts	stdin	urandom
mqueue	random	stdout	zero

```
imsrv01@imsrv01-Lenovo-ideapad-700-15ISK:~/CKA/cka/security$ k exec -it privileged-pod -- ls /dev
```

autofs	stdin	tty48
bsg	stdout	tty49
btrfs-control	termination-log	tty5
core	tty	tty50
cpu	tty0	tty51
cpu_dma_latency	tty1	tty52
fd	tty10	tty53
full	tty11	tty54
fuse	tty12	tty55
hpet	tty13	tty56
hwrng	tty14	tty57
input	tty15	tty58
kmsg	tty16	tty59
loop-control	tty17	tty6
loop0	tty18	tty60
loop1	tty19	tty61
loop2	tty2	tty62
loop3	tty20	tty63
loop4	tty21	tty7
loop5	tty22	tty8
loop6	tty23	tty9
loop7	tty24	ttyS0
mapper	tty25	ttyS1
mem	tty26	ttyS2
memory_bandwidth	tty27	ttyS3
mqueue	tty28	urandom
net	tty29	usbmon0
network_latency	tty3	vboxguest
network_throughput	tty30	vboxuser
null	tty31	vcs
nvr	tty32	vcs1
port	tty33	vcs2
ptmx	tty34	vcs3
pts	tty35	vcs4
random	tty36	vcs5
rfs	tty37	vcs6
rtc0	tty38	vcsa
sda	tty39	vcsa1
sda1	tty4	vcsa2
sda2	tty40	vcsa3
sg0	tty41	vcsa4
sg1	tty42	vcsa5

shm	tty43	vcsa6
snapshot	tty44	vga_arbiter
snd	tty45	vhost-vsock
sr0	tty46	zero
stderr	tty47	

## Linux Capabilities

- The capabilities to add/drop when running containers. Defaults to the default set of capabilities granted by the container runtime.
- With [Linux capabilities](#), you can grant certain privileges to a process without granting all the privileges of the root user.
- Find capabilities  

```
cd /proc/1 (1 - proces id)
```

```
cat status
```

```
securityContext:
  capabilities:
    add:
      - SYS_TIME
```

```
securityContext:
  capabilities:
    drop:
      - CHOWN
```

## ReadOnlyRootFilesystem

- Whether this container has a read-only root filesystem. Default is false

```
securityContext:
  readOnlyRootFilesystem: true
volumeMounts:
  - name: my-volume
    mountPath: /volume
  readOnly: false
```

```
Cannot write to root file system but can write to mounted volume..
```

## SELinux

```
...
securityContext:
  seLinuxOptions:
    level: "s0:c123,c456"
```

### AppArmor

AppArmor is a Linux kernel security module that supplements the standard Linux user and group based permissions to confine programs to a limited set of resources. AppArmor can be configured for any application to reduce its potential attack surface and provide greater in-depth defense. It is configured through profiles tuned to whitelist the access needed by a specific program or container, such as Linux capabilities, network access, file permissions, etc. Each profile can be run in either *enforcing* mode, which blocks access to disallowed resources, or *complain* mode, which only reports violations.

## Secure Images

The default pull policy is `IfNotPresent` which causes the Kubelet to skip pulling an image if it already exists. If you would like to always force a pull, you can do one of the following:

- set the `imagePullPolicy` of the container to `Always`.
- omit the `imagePullPolicy` and use `:latest` as the tag for the image to use.
- omit the `imagePullPolicy` and the tag for the image to use.
- enable the [AlwaysPullImages](#) admission controller.

Note that you should avoid using `:latest` tag

## Access private repositories

Example : Accessing docker/ecr repo using service account with secrets - `imagePullSecrets`



# Authorization

## Modes

**Node** - A special-purpose authorization mode that grants permissions to kubelets based on the pods they are scheduled to run

The Node authorizer allows a kubelet to perform API operations. This includes:

Read operations:

- services
- endpoints
- nodes
- pods
- secrets, configmaps, persistent volume claims and persistent volumes related to pods bound to the kubelet's node

Write operations:

- nodes and node status (enable the NodeRestriction admission plugin to limit a kubelet to modify its own node)
- pods and pod status (enable the NodeRestriction admission plugin to limit a kubelet to modify pods bound to itself)
- events

Auth-related operations:

- read/write access to the certificationsigningrequests API for TLS bootstrapping
- the ability to create tokenreviews and subjectaccessreviews for delegated authentication/authorization checks

## RBAC -

- Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within an enterprise
- To enable RBAC, start the apiserver with `--authorization-mode=RBAC`

## Webhook -

- A WebHook is an HTTP callback: an HTTP POST that occurs when something happens; a simple event-notification via HTTP POST. A web application implementing WebHooks will POST a message to a URL when certain things happen

**ABAC** - Attribute-based access control (ABAC) defines an access control paradigm whereby access rights are granted to users through the use of policies which combine attributes together. The policies can use any type of attributes (user attributes, resource attributes, object, environment attributes, etc).

1. Alice can do anything to all resources:  

```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"user": "alice", "namespace": "*", "resource": "*", "apiGroup": "*"}}
```
2. The Kubelet can read any pods:  

```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"user": "kubelet", "namespace": "*", "resource": "pods", "readOnly": true}}
```
3. The Kubelet can read and write events:  

```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"user": "kubelet", "namespace": "*", "resource": "events"}}
```
4. Bob can just read pods in namespace "projectCaribou":  

```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"user": "bob", "namespace": "projectCaribou", "resource": "pods", "readOnly": true}}
```

Anyone can make read-only requests to all non-resource paths:

- ```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"group": "system:authenticated", "readOnly": true, "nonResourcePath": "*"}}
```
5. 

```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"group": "system:unauthenticated", "readOnly": true, "nonResourcePath":
```

For example, if you wanted to grant the default service account (in the kube-system namespace) full privilege to the API using ABAC, you would add this line to your policy file:

```
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec": {"user": "system:serviceaccount:kube-system:default", "namespace": "*", "resource": "*", "apiGroup": "*"}}
```

- `--authorization-mode=ABAC` Attribute-Based Access Control (ABAC) mode allows you to configure policies using local files.
- `--authorization-mode=RBAC` Role-based access control (RBAC) mode allows you to create and store policies using the Kubernetes API.
- `--authorization-mode=Webhook` WebHook is an HTTP callback mode that allows you to manage authorization using a remote REST endpoint.
- `--authorization-mode=Node` Node authorization is a special-purpose authorization mode that specifically authorizes API requests made by kubelets.
- `--authorization-mode=AlwaysDeny` This flag blocks all requests. Use this flag only for testing.
- `--authorization-mode=AlwaysAllow` This flag allows all requests. Use this flag only if you do not require authorization for your API requests.

## SelfSubjectAccessReview

`kubectl auth can-i`

- for quickly querying the API authorization layer.
- uses the `SelfSubjectAccessReview` API to determine if the current user can perform a given action, and works regardless of the authorization mode used

```
kubectl auth can-i create deployments --namespace dev
yes
```

```
kubectl auth can-i create deployments --namespace prod
no
```

Others:

- **SubjectAccessReview** - Access review for any user, not just the current one. Useful for delegating authorization decisions to the API server. For example, the kubelet and extension API servers use this to determine user access to their own APIs.
- **LocalSubjectAccessReview** - Like `SubjectAccessReview` but restricted to a specific namespace.
- **SelfSubjectRulesReview** - A review which returns the set of actions a user can perform within a namespace. Useful for users to quickly summarize their own access, or for UIs to hide/show actions.

## Userimpersonation

```
$kubectl auth can-i create deployment -n kube-system --as default
```

no - no RBAC policy matched

## Network Policies

## Creating TLS certificates

## Admission Controller

### **AlwaysPullImages**

- This admission controller modifies every new Pod to force the image pull policy to Always.
- This is useful in a multitenant cluster so that users can be assured that their private images can only be used by those who have the credentials to pull them.
- When this admission controller is enabled, images are always pulled prior to starting containers, which means valid credentials are required